

# Deploying LLMs on Low-Power Edge Devices: Techniques, Algorithms, and RISC-V Optimizations

Author: Generated by GPT-5 Thinking mini  
Date: September 21, 2025

**Abstract:** This document provides an in-depth, practical guide to deploying large language models (LLMs) on constrained edge devices, with a focus on techniques that reduce inference time, memory footprint, and power consumption. It covers model compression (quantization, pruning, distillation), efficient transformer variants, runtime/kernel optimizations, and concrete RISC-V-specific hardware and ISA strategies. Additionally, it discusses tooling, benchmarking, security, and a step-by-step implementation plan for engineers targeting IoT-class hardware.

## 1. Introduction and problem statement

Introduction and Problem Statement Running state-of-the-art LLMs on cloud servers yields excellent performance but depends on network connectivity, latency, data privacy, and operational costs. For many applications — industrial sensors, wearable devices, remote monitoring, and privacy-sensitive workloads — on-device inference is desirable or required. Edge devices impose strict constraints: limited RAM (often 64KB to a few MB for microcontrollers, up to a few GB for higher-end SoCs), tight energy budgets, lower compute (small integer ALUs, limited or no floating-point units), and reduced storage. Achieving LLM-like capabilities under these constraints requires rethinking models, runtimes, and hardware support. This document frames the problem: deliver useful language capabilities (inference, classification, retrieval-augmented responses) with acceptable latency (tens to hundreds of milliseconds where possible), minimal power draw, and a small memory footprint while maintaining privacy and robustness. The remainder of the PDF outlines approaches at the algorithmic, system, and ISA levels to meet this challenge.

## 2. Background — LLMs, architecture, and compute profile

Background — LLMs, architecture, and compute profile Large language models are built from transformer blocks: multi-head self-attention, feed-forward networks (FFNs), normalization layers, and embeddings. The dominant computational kernels are matrix multiplications (GEMM), softmax and layer normalization. These kernels are heavy in memory bandwidth and require efficient implementation for quantized data types. Key compute characteristics:

- Dominant GEMM workloads: attention and FFN layers.
- Memory-bound vs compute-bound: smaller models tend to be compute-bound on tiny devices, while larger layers become memory-bandwidth-bound.
- Sensitivity to numerical precision: transformers often tolerate reduced precision (8-bit, 4-bit, and in some cases 2-bit quantization) when properly calibrated with quantization-aware training or post-training techniques. Understanding these traits informs which algorithmic and hardware optimizations will yield the largest real-world gains.

### **3. Model compression — pruning, quantization, and distillation**

Model compression techniques

- 1) Quantization Quantization maps high-precision floating values to lower bit-width integer representations. Common strategies:
  - Post-Training Quantization (PTQ): simple and fast, but may degrade accuracy for aggressive bit-widths.
  - Quantization-Aware Training (QAT): simulates quantization during training to retain accuracy.
  - Mixed precision: keeping sensitive layers (e.g., layernorm, embedding tables) at higher precision while quantizing the heavy GEMM layers to 8/4/2 bits.
- Recent techniques: symmetric vs asymmetric quantization, per-channel scale factors for weights, group-wise quantization, and learned step size quantization (LSQ). For LLMs, per-channel weight quantization for linear layers usually provides a better accuracy/size tradeoff.
- 2) Pruning Pruning removes redundant weights or entire neurons. Methods include magnitude pruning, structured pruning (removing attention heads or MLP dimensions), and lottery-ticket style rewinding. Structured pruning is attractive for edge devices because it yields dense smaller matrices friendly to optimized kernels.
- 3) Distillation Knowledge distillation trains a smaller student model to emulate a larger teacher. Distillation can be task-specific (fine-tune a distilled model for a subset of tasks) or general (distil large LLM behavior). Tiny student models (millions to hundreds of millions of parameters) can capture much of the teacher's performance for narrow tasks. Combining techniques (quantize after distillation, prune then fine-tune with QAT) typically yields the best results.

## 4. Efficient architectures and algorithmic changes

Efficient architectures and algorithmic changes Redesigning the model architecture for efficiency can pay large dividends:

- Lightweight transformer variants: Linformer, Performer, Longformer, Reformer — each reduces attention complexity or approximates it to reduce compute and memory.
- Funnel-Transformers and grouped attention: compress sequence length or compute attention selectively to lower cost.
- Sparse Mixture-of-Experts (MoE): activate only a subset of parameters per input; challenges include routing overhead and load balancing.
- Low-rank or factorized FFNs: approximate large dense FFNs with low-rank decompositions.
- Token reduction and early exit mechanisms: reduce tokens dynamically based on attention scores or confidence thresholds. For edge scenarios, hybrid solutions work well: small core transformer + retrieval-augmented generation (RAG) where a compact local memory or a nearby gateway supplies extra context.

## 5. RISC-V specific optimizations and custom instructions

RISC-V specific optimizations and custom instructions RISC-V is attractive for edge deployment due to its extensible ISA and growing ecosystem. Important vectors for optimization: 1) Vector Extension (RVV) RVV provides a scalable vector compute model enabling SIMD-like operations with flexible vector lengths. Implementing GEMM, convolution, and attention kernels with RVV yields dramatic speedups compared to scalar implementations. Key ideas: - Tile and strip-mining strategies to fit vector register windows. - Use of vector reductions and widening operations for accumulation to preserve precision. 2) Custom Instructions and Accelerator Coprocessors Designing custom instructions for common kernels (e.g., fused multiply-add for int8, matrix-vector kernels, specialized softmax approximations) reduces instruction overhead. In conjunction, lightweight accelerators (small matrix-multiply units, systolic arrays, or TPU-like MAC arrays) can sit alongside the core. 3) Memory hierarchy and DMA Use on-chip scratchpads and DMA to stream tiles from flash/DDR into local SRAM. Explicit software-managed scratchpads outperform caching in predictable access patterns typical of neural kernels. 4) Fixed-point arithmetic and mixed precision Leverage integer-only math where possible. For RISC-V cores lacking FPUs, integer arithmetic with scale factors and block-floating approaches maintain dynamic range while enabling fast integer arithmetic. 5) Instruction scheduling and loop unrolling Optimize loops to minimize load/store stalls, prefetch upcoming tiles, and align data to allow vector loads/stores. Carefully crafted assembly kernels or compiler intrinsics are essential. Combining RVV, custom instructions, and a small accelerator yields the best energy-to-inference tradeoff on RISC-V platforms.

## 6. Runtime and kernel optimizations

Runtime and kernel optimizations

- 1) Memory planning Compute memory requirements for activations, weights, and temporary buffers, and schedule operators to reuse buffers. Peak memory often occurs during attention due to storing QKV or past key-values; use recomputation/trading compute for memory to reduce peak usage.
- 2) Operator fusion Fuse common sequences (e.g., linear -> GELU -> dropout -> add) to eliminate intermediate buffers and cache thrash. Fusion reduces overhead and increases locality.
- 3) Low-level kernels Write hand-optimized kernels for int8/4 operations: blocked GEMM, optimized im2col where applicable, and fused attention kernels that compute Q, K, V and attention in a streaming fashion. Use register blocking and minimize memory spills.
- 4) Runtime adaptivity Adjust batch sizes, sequence lengths, and precision at runtime based on resource availability. For example, when battery is low, switch to int4 or a smaller model.
- 5) Cross-platform toolchains Toolchains like TVM, Glow, XLA (for edge), and Vela (for RISC-V) allow ahead-of-time compilation and optimization. Use autotuning to find kernel parameters (block sizes, vector lengths) that suit the hardware.

## 7. Retrieval, context management, and latency reduction

Retrieval, context management, and latency reduction Full LLM decoding for long contexts is expensive. Strategies to reduce latency:

- Retrieval-Augmented Generation (RAG): store local embeddings (or a compressed index) and retrieve relevant chunks to present to a small local model, reducing the need for large model computation.
- Compressed and hierarchical memory: use a short-term token cache for recent context and a compressed long-term memory with downsampled representations.
- Token-level early exit: implement confidence-based early stopping during autoregressive decoding for tokens with high certainty.
- Caching past key-values: for streaming inputs, cache K/V states efficiently in compressed/int8 form to avoid recomputation. These techniques avoid running the full model on every input and greatly reduce average latency and energy usage.

## 8. Benchmarking, metrics, and evaluation

Benchmarking, metrics, and evaluation When optimizing for edge, track metrics beyond raw accuracy:

- Latency (p99 and median), throughput (tokens/sec), and jitter.
- Energy per inference (mJ), measured with power instrumentation across typical workloads.
- Memory footprint (peak RAM), flash storage, and binary size.
- Quality metrics: perplexity, task-specific accuracy, and human evaluations for generation quality.
- Robustness: performance under varied thermal and power states, and sensitivity to quantization.

Design benchmark suites with representative workloads (short queries, streaming inputs, sensor-triggered prompts). Use end-to-end measurements including sensor I/O, preprocessing, inference, and postprocessing to avoid overly optimistic microbenchmarking.

## **9. Security, privacy, and deployment considerations**

Security, privacy, and deployment considerations Running models on-device enhances privacy, but new risks arise:

- Model extraction risk: adversaries can probe models to reconstruct parameters. Use access controls and rate limits.
- Data leakage: ensure local logs and caches are encrypted at rest and erased securely.
- Adversarial inputs: deploy input sanitization, anomaly detectors, and robust training techniques.
- Secure boot and signed firmware: ensure only authenticated runtime images execute on devices.
- Update strategies: support secure model updates (delta-updates, signed packages) and graceful rollback in case of failure. For regulatory compliance (e.g., data protection laws), document what data stays on-device and how long logs are retained. Consider hybrid architectures where private data is processed locally while non-sensitive heavy tasks are offloaded.

## 10. Implementation guide, tools, and future directions

Implementation guide, tools, and future directions Practical steps to go from prototype to production: 1) Choose model family: start with a small transformer (e.g., distilled models or 10M-100M parameter families) and iterate. 2) Prototype with quantization-aware training and distillation. Evaluate on your target tasks. 3) Benchmark on target hardware (use cycle-accurate simulators or dev boards). Measure energy and latency. 4) Implement optimized kernels using RVV intrinsics or assembler; consider a tiny accelerator for GEMM. 5) Integrate runtime optimizations: memory planner, fused operators, and adaptive precision switching. 6) Harden security: secure boot, signed updates, and encrypted storage. 7) Automate CI/CD for model builds and firmware images. Tooling: TVM, ONNX Runtime (with micro runtime), TensorFlow Lite Micro, XNNPACK, Apache TVM with RISC-V backends, and vendor SDKs. For RISC-V, look into RVV-enabled toolchains and frameworks providing vectorized kernels. Future directions: ultra-low-bit quantization (1–2 bit), learned compression codecs for activations, co-designed hardware-software stacks with tiny accelerators, and federated distillation where devices collaboratively improve a shared compact model without sharing raw data. Conclusion: Deploying LLM-like capabilities to edge devices is attainable by combining model compression, architecture redesign, runtime and ISA-level optimizations, and careful benchmarking. RISC-V's modular ISA and vector extensions make it a strong platform for future low-power LLM deployments.

References and further reading (select): - Papers on quantization, distillation, Linformer, Performer, RVV specs, TVM documentation, and TinyML literature.